TRENDYTECH
UPLIFT YOUR CAREER!

Newly Launched Premium Program

The Ultimate Big Data Masters Program
(Cloud Focused)

By

Sumit Mittal

CURRICULUM

Module 1 - Big Data Fundamentals

- Introduction to Big Data
- How Big Data Works
- Practice Environment
- Brief Introduction to Distributed System Architecture
- Introduction to Hadoop and its Ecosystem Tools
- Basics of Distributed Storage - HDFS Architecture
- Linux Commands
- HDFS Commands & How it Works
- Introduction to Data Lake Storage - Blob & ADLS Gen 2
- Big Data - The Big Picture with Real-Time Example

Module 2 - Distributed Processing with Pyspark

- Distributed Processing Fundamentals
- Knowing Apache Spark
- Spark Development Environments - OnPremise | OnCloud
- Understanding Spark Cluster & Cluster Modes
- Apache Spark In-Depth with Real-Time Example
- How Spark Executes Program on the Cluster
- Stages in Spark
- Understanding Spark Transformations & Actions
    - Lazy Evaluation
    - Narrow Vs Wide Transformations
- Accumulators & Broadcast Variables
- Repartition Vs Coalesce
- Data Caching
    - Spark Storage Levels
    - Cache Vs Persist

- Spark Optimization Techniques In-depth
- Internals of File Formats - Parquet | ORC | Avro
- Compression Techniques
- Introduction to Spark Data Frames
  - Creating Spark Data frame
  - Data frame Transformations and Actions
  - Querying Spark Data frame More Data frame Transformations
- Introduction to SparkSQL
- Understanding Cluster Configurations
- How to Submit Spark Job
- Scheduling and Running Spark Jobs
- Spark Advance Optimizations - Sort Vs Hash Aggregate
- Spark Catalyst Optimizer
- Learning Hive
- Spark-Hive Integration
- Implement Your First Batch Processing Project with Pyspark

Module 3 - Azure Databricks

- Introducing Azure Databricks
- Microsoft Azure Services and Portal Overview
- What is Databricks & Why Databricks Databricks
- Pricing - Infrastructure and Software Charges
- Different Cloud Providers offering Databricks
- Databricks Features
- Databricks Community Edition
- 3 ways to Create Cluster
- All Purpose Cluster
- Job Cluster
- Cluster Pool
- When to use the Different Cluster Modes
- Databricks Benefits
- Different optimized Cluster types - Memory Optimized, Storage Optimized, Compute Optimized, General Purpose, GPU Accelerated

- Databricks File System (DBFS)
- Databricks Architecture - Control and Data Plane
- DBFS in detail
- Object Store - Blob, Datalake Gen2
- Filesystem utility- dbutils
- Data Utility & Notebook Utility & Widgets Utility
- Parameter passing from one Notebook to another
- Mount Point - How to create Mount Point
- Databricks Workspace
- Databricks CLI
- Ways to access Storage Account
- Access Key | Account Key
- SAS Key & Service Principal
- Secret Scope - Azure Key vault Backed Secret Scope & Databricks Backed Secret Scope

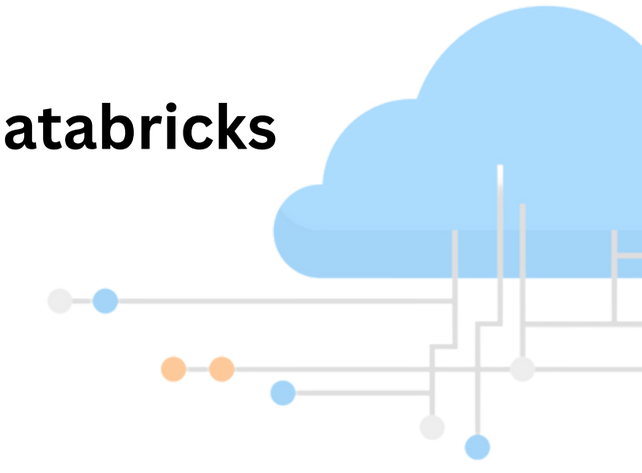

- **Delta Lake**
- **Delta Table Creation**
- **Lakehouse Architecture**
- **Azure Delta Engine Optimizations**
- **Delta Architecture - Medallian Architecture**
- **Cluster Creation**
- **Autoloader**
- **Delta Live Table**
- **Unity Catalog**

Module 4 - Azure DataFactory

- Azure Data Factory Introduction
- Data Transfer (Source to Sink)
- Data Transformation - Data Flow
- Workflow Orchestration
- Data Transfer from RDBMS to ADLS Gen2
- Azure SQL Databses
- Data Transfer from Azure SQL to ADLS Gen2
- Author, Monitor & Manage
- Data Integration Service (ADF)
- Usecases where ADF can be used
- Data Ingestion
- Data Transformation
- Data Orchestration
- Data Flow Mapping
- Data transfer from external URL to ADLS - Usecase

- Linked Services for Source and Sink
- Select Transformation
- ADF Primary Usage
- Tansfer data from Blob to Datalake - Usecase
- Blob Connector
- Http Connector
- Datalake Instance
- Data Factory Instance
- Linked Service Creation - Blob & Datalake
- Dataset for Blob and Datalake
- Complete Pipeline setup
- Key Vault & Scheduled Triggers
- Tumbling Window Triggers
- Storage & Custom Events
- Trigger Pipeline on Custom Event - Usecase
- Data Ingestion from 2 Sources (Blob & Amazon S3) to ADLS Gen2
- Building a Complete Pipeline Using DataBricks & DataFactory

Module 5 - Interview Readiness

- **Data Modeling**
  - **Fact & Dimension Tables**
  - **Data Models - Star Vs Snowflake**
- **System Design**
- **CICD - Git**
- **Interview Preparation Tips**
- **Interview Questions**
- **Guidance for Resume Preparation**
- **How to Handle Managerial Round Questions**

Module 6 - Streaming

- Structured Streaming In-depth
- Benefits of Spark Structured Streaming
- Types of Data Sources
- Streaming Joins
- Streaming Dataframe
- Introduction to Kafka - Streaming Platform
- Kafka Architecture
- Installing Multi-Node Kafka Cluster
- Writing Kafka Producer and Consumer
- Scaling up the Kafka Cluster
- Integrating Kafka with Spark Structured Streaming
- Building Streaming Pipeline (Structured Streaming with Kafka)

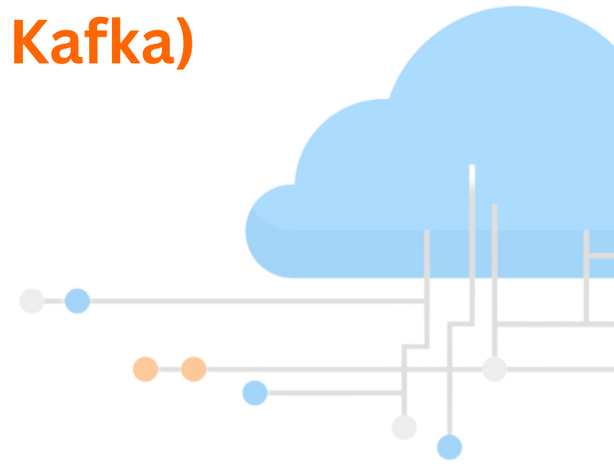

Module 7 - More on Cloud Services

- Azure Synapse
- Azure CosmosDB
- Azure HDInsights
- Azure Logic App
- Azure Event Hub

- AWS EMR (Elastic MapReduce)
- Launch EMR Cluster Using Advanced Options
- Types of EC2 Instances
- AWS S3
- AWS Athena
- AWS Glue - Data Catalog | Crawlers
- AWS Redshift
- End-to-End Real-time Project on Cloud using other cloud services

Additional Modules

(To Crack Top Product Based Companies)

- **Data Structures and Algorithms for Data Engineers - 20 hours**
- **Python for Data Engineers - 10 hours**

TRENDY TECH
UPLIFT YOUR CAREER!

Contact

hello@trendytech.in

https://trendytech.in/